

Capstone Project Report

June 30, 2020

Report in Jupyter notebook exported as pdf. Please visit <https://github.com/abiswas20> for additional examples of personal data science projects.

```
[2]: from IPython.display import Image
```

1 Identifying Potential Toronto Neighborhood(s) To Open A Profitable Coffee Shop

1.1 The Problem

In this hypothetical situation, our client, who is an entrepreneur, wants to open an independent coffee shop in Toronto. My job, as a data scientist working on the project, is to use data science and analysis to recommend a list of such possible neighborhoods.

1.2 Introduction

Toronto is the largest city and the financial capital of Canada. And as one would expect, its market for coffee shops, for the most part, is already being served by existing businesses. With that in perspective, we give special attention to the following factors to search for opportunities remaining in this highly competitive market:

1. Population: The neighborhood has to have a large enough population for a new coffee shop to find its own patrons.
2. Average Income: After accounting for employees, rent and utilities the price of a cup of coffee will have to be on the higher end for a business to survive. Being in a neighborhood with average income in the top 25% percent would offset for such pricing and would be ideal.
3. Walkability: New coffee shops usually do better when there is foot traffic and people take the risk of exploring new venues. A higher Walk Score would make for a better potential candidate.
4. Business atmosphere: We want to see a lot of already successful businesses and low debt risk score in the neighborhood, general indicators of a good business environment. Moreover, the existence of multiple successful coffee shops would be a huge plus.
5. Parks and Playgrounds: Such public spaces for recreation invites lots of people, including new faces, everyday. This is especially true in large cities.

We want to be in a cluster which has high scores is as many of the five factors above as possible.

1.3 Data

Sources of data used in the analysis are listed below: 1. Geographical Coordinates - Foursquare API[1] and OpenStreetMap Nominatim package[2]

2. Population, Average income - Wikipedia[3]
3. Walkability - Wellbeing Toronto Civics Equity Indicators (Toronto Open Data)[4]
4. Businesses, Debt Risk Score - Wellbeing Toronto Economics (Toronto Open Data)[4]
5. Venues including Coffee shops, Parks and Playgrounds - Foursquare API[1]

Data on population, average income, walkability, neighborhood businesses including coffee shops and recreational spaces were analyzed to evaluate the five factors of each neighborhood, as highlighted in the introduction. They were combined together into a single dataframe as all of them add to the positive outlook of a neighborhood. The data was then normalized to ensure they all have equal weight in any further analysis.

1.4 Methodology

The goal of this project is to create a short list of 1-5 most prospective neighborhoods. The introduction section explains specific reasons behind why certain features were selected. We are looking for a location in an affluent neighborhood or close to it, with a variety of successful businesses around, including other highly rated coffee shops. A walkable neighborhood opens up the possibility of people taking a chance on a new venue. Recreational spaces also draw significant crowd in big cities, bringing new people to the area who are also potential customers.

We are looking to create a short list, which makes the problem well suited for K Means Clustering. We start out by scraping the table from the Wikipedia article on demographics of Toronto neighborhoods [3]. The table is cleaned up and only neighborhood name, population and average income are included in the `df_demographics` dataframe.

Data on the Walk Score of each neighborhood came from the 'Wellbeing Toronto Civics Equity Indicators' dataset obtained through Toronto Open Data initiative[4]. I saved the Walk Score data in the `df_walkability` dataframe. The data for the third dataframe, `df_economics`, came from 'Wellbeing Toronto Dataset', again obtained through Toronto Open Data initiative[4]. As the name suggests, the dataset contains economic indicators of Toronto neighborhoods. Only data on number of businesses and debt risk score are saved in the `df_economics` dataframe.

The three dataframes `df_demographics`, `df_walkability` and `df_economics` are combined into a single dataframe `df_combined`. I clustered them into 8 clusters using K means clustering algorithm and chose only two clusters whose centroids were farthest from the origin in terms of Euclidean distance. Each of the two clusters had one neighborhood each, giving us a short list of only two neighborhoods.

Finally, it's important to note that I could find complete data on only 74 Toronto neighborhoods. The project focuses on those 74 neighborhoods.

1.5 Results

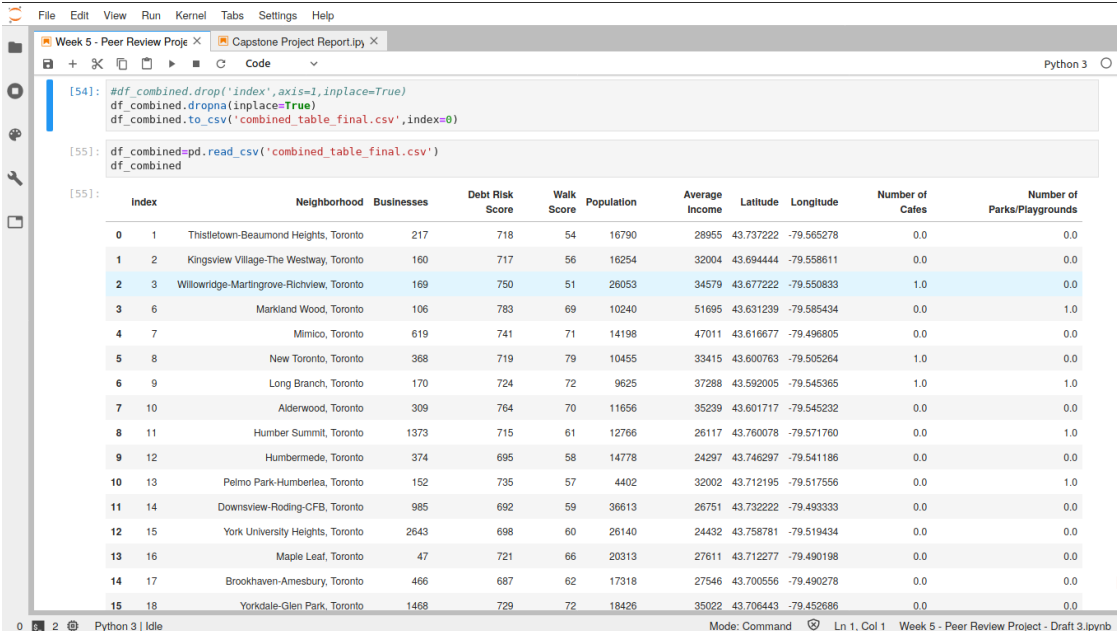
As mentioned in an earlier section three dataframes were merged to create a single dataframe `df_combined`:

1. `df_demographics`;
2. `df_walkability`; and,
3. `df_economics`.

Futhermore, data from Foursquare on the number of cafes and playgrounds/parks within top 100 venues in the neighborhood were added to the dataframe. Only those 74 neighborhoods that had complete information for all the fields in `df_combined` were chosen for further analysis. The neighborhoods were plotted on a map of Toronto to illustrate their spatial distribution (see screenshots below). Geographical coordinates were removed and the data was normalized before performing K means clustering.

```
[3]: Image(filename="Screenshot from 2020-05-24 12-05-13.png")
```

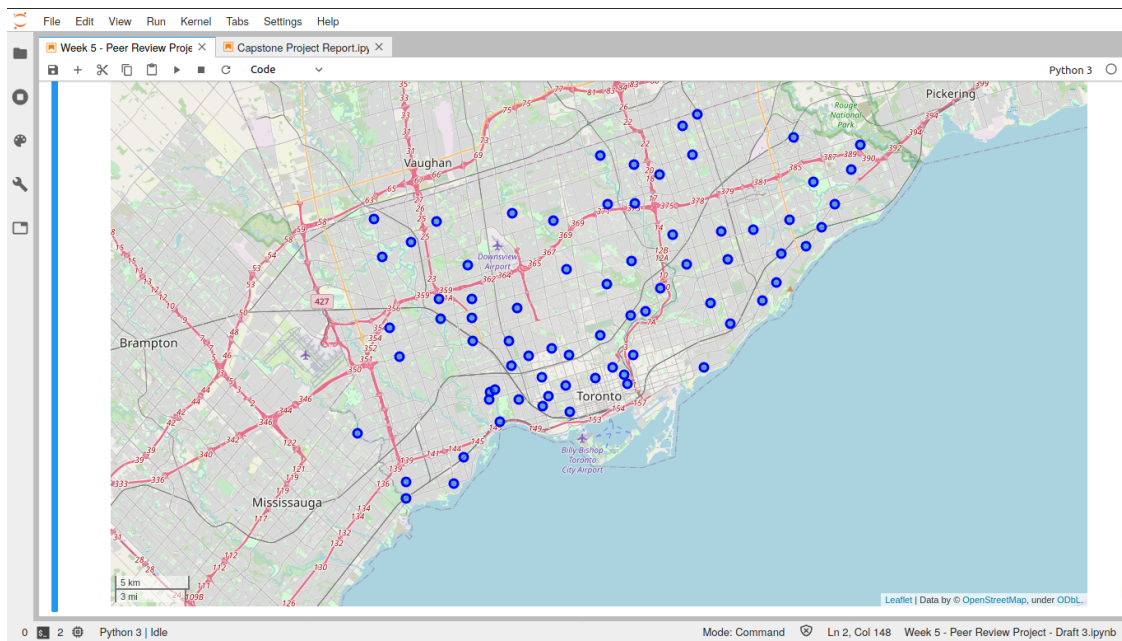
[3]:



	Index	Neighborhood	Businesses	Debt Risk Score	Walk Score	Population	Average Income	Latitude	Longitude	Number of Cafes	Number of Parks/Playgrounds
0	1	Thistletown-Beaumont Heights, Toronto	217	718	54	16790	28955	43.737222	-79.565278	0.0	0.0
1	2	Kingsview Village-The Westway, Toronto	160	717	56	16254	32004	43.694444	-79.558611	0.0	0.0
2	3	Willowridge-Martingrove-Richview, Toronto	169	750	51	26053	34579	43.677222	-79.550833	1.0	0.0
3	6	Markland Wood, Toronto	106	783	69	10240	51695	43.631239	-79.585434	0.0	1.0
4	7	Mimico, Toronto	619	741	71	14198	47011	43.616677	-79.496805	0.0	0.0
5	8	New Toronto, Toronto	368	719	79	10455	33415	43.600763	-79.505264	1.0	0.0
6	9	Long Branch, Toronto	170	724	72	9625	37288	43.592005	-79.545365	1.0	1.0
7	10	Alderwood, Toronto	309	764	70	11656	35239	43.601717	-79.545232	0.0	0.0
8	11	Humber Summit, Toronto	1373	715	61	12766	26117	43.760078	-79.571760	0.0	1.0
9	12	Humbermede, Toronto	374	695	58	14778	24297	43.746297	-79.541186	0.0	0.0
10	13	Pelmo Park-Humberlea, Toronto	152	735	57	4402	32002	43.712195	-79.517556	0.0	1.0
11	14	Downsview-Roding-CFB, Toronto	985	692	59	36613	26751	43.732222	-79.493333	0.0	0.0
12	15	York University Heights, Toronto	2643	698	60	26140	24432	43.758781	-79.519434	0.0	0.0
13	16	Maple Leaf, Toronto	47	721	66	20313	27611	43.712277	-79.490198	0.0	0.0
14	17	Brookhaven-Amesbury, Toronto	466	687	62	17318	27546	43.700556	-79.490278	0.0	0.0
15	18	Yorkdale-Glen Park, Toronto	1468	729	72	18426	35022	43.706443	-79.452686	0.0	0.0

```
[4]: Image(filename='Screenshot from 2020-05-24 12-05-39.png')
```

[4]:



The clustering method from SciKit Learn package was used to carry out K means clustering. It was set to produce 8 clusters where centroid of each cluster is defined by corresponding values of the 7 features. The coordinates of the cluster centroids and their Euclidean distance from the origin are shown in the `df_centroids` dataframe below. The centroids of clusters 5 and 6 stand out for their large Euclidean distance. Both are approximately twice farther away from the origin than the centroid of cluster 0, which comes in a distant third.

The `clustering_data` dataframe attaches labels to each neighborhood specifying the cluster to which they belong.

```
[5]: Image(filename='Screenshot from 2020-05-25 13-28-55.png')
```

```
[5]:
```

```

File Edit View Run Kernel Tabs Settings Help
Week 5 - Peer Review Proj... Capstone Project Report.ipynb Python 3
[168]: df_centroids['Euclidean Distance'] = ''
        from scipy.spatial import distance
        #from math import dist
        i=0
        while i<len(df_centroids.index):
            centroid=(df_centroids.iloc[i,0], df_centroids.iloc[i,1], df_centroids.iloc[i,2], df_centroids.iloc[i,3], df_centroids.iloc[i,4], df_centroids.iloc[i,5], df_centroids.iloc[i,6], df_centroids.iloc[i,7])
            df_centroids.iloc[i,7]=distance.euclidean((0, 0, 0, 0, 0, 0, 0, 0), centroid)
            i+=1
        df_centroids

[168]:
      Businesses  Debt Risk Score  Walk Score  Population  Average Income  Number of Cafes  Number of Parks/Playgrounds  Euclidean Distance
Centroid 0   -0.113481      0.183924   1.366639   -0.933358      0.120070      2.103009      -0.304421      2.70468
Centroid 1   -0.324407      1.260738   -0.108139   0.016303      0.391602     -0.525446      0.019245      1.46167
Centroid 2   -0.188619     -0.464789   -0.681680   0.057142     -0.337416     -0.400821     -0.514804      1.12209
Centroid 3    0.942822     -0.172983   -0.575749   2.078678     -0.460172     -0.525446     -0.336788      2.48445
Centroid 4   -0.149553      0.790275   1.320414   -0.280708      0.786692      0.139221      3.104861      3.57038
Centroid 5    0.266437      1.421529   1.150925   -0.914492      6.890640     -0.525446     -0.692820      7.24498
Centroid 6    6.091906      0.635439   2.422096   -1.193605      0.067392      1.468555      0.731310      6.89219
Centroid 7   -0.261284     -0.963540   -0.007253   -0.417660     -0.361829     -0.276196      0.731310      1.38318

Let's save a copy of df_centroids before moving on.

[169]: df_centroids.to_csv('centroids.csv')


[184]: df_centroids=pd.read_csv('centroids.csv')

[191]: df_centroids.index=df_centroids['Unnamed: 0']
        df_centroids.drop('Unnamed: 0',axis=1,inplace=True)
        del df_centroids.index.name
        df_centroids

```

[6]: Image(filename='Selection_001.png')

[6]:



[7]: Image(filename='Screenshot from 2020-05-25 13-29-23.png')

[7]:

	Cluster Label	Neighborhood	Businesses	Debt Risk Score	Walk Score	Population	Average Income	Number of Cafes	Number of Parks/Playgrounds
0	2	Thistletown-Beaumont Heights, Toronto	217	718	54	16790	28955	0.0	0.0
1	2	Kingsview Village-The Westway, Toronto	160	717	56	16254	32004	0.0	0.0
2	2	Willowridge-Martingrove-Richview, Toronto	169	750	51	26053	34579	1.0	0.0
3	1	Markland Wood, Toronto	106	783	69	10240	51695	0.0	1.0
4	2	Mimico, Toronto	619	741	71	14198	47011	1.0	0.0
5	0	New Toronto, Toronto	368	719	79	10455	33415	3.0	0.0
6	7	Long Branch, Toronto	170	724	72	9625	37288	1.0	1.0
7	1	Alderwood, Toronto	309	764	70	11656	35239	0.0	0.0
8	7	Humber Summit, Toronto	1373	715	61	12766	26117	0.0	1.0
9	2	Humbermede, Toronto	374	695	58	14778	24297	0.0	0.0
10	7	Peimo Park-Humberlea, Toronto	152	735	57	4402	32002	0.0	1.0
11	3	Downsview-Roding-CFB, Toronto	985	692	59	36613	26751	0.0	0.0
12	3	York University Heights, Toronto	2643	698	60	26140	24432	0.0	0.0
13	2	Maple Leaf, Toronto	47	721	66	20313	27611	0.0	0.0
14	2	Brookhaven-Amesbury, Toronto	466	687	62	17318	27546	0.0	0.0
15	2	Yorkdale-Glen Park, Toronto	1468	729	72	18426	35022	0.0	0.0
16	2	Bathurst Manor, Toronto	244	745	61	14945	34169	0.0	1.0
17	2	Westminster-Branson, Toronto	83	730	61	16386	27826	0.0	0.0
18	1	Lansing-Westgate, Toronto	488	758	77	10052	46631	0.0	0.0
19	1	Bedford Park-Nortown, Toronto	676	776	73	13749	80827	0.0	0.0
20	1	Bridle Path-Sunnybrook-York Mills, Toronto	58	788	58	17564	92099	0.0	1.0
21	1	Banbury-Don Mills, Toronto	834	776	67	21372	47515	0.0	1.0

Clusters 5 and cluster 6 have only one neighborhood each. The two clusters were combined to create the Preferred_Neighborhood dataframe. As the name suggests, the two neighborhoods listed in the dataframe: Bay St Corridor and Rosedale-Moore Park, have the highest prospect at absorbing a new coffee shop and making it a success.

[8]: `Image(filename='Screenshot from 2020-05-25 13-30-12.png')`

[8]:

```
[195]: Preferred_Neighborhoods=pd.concat([cluster_5,cluster_6],ignore_index=True)
Preferred_Neighborhoods
```

	Cluster Label	Neighborhood	Businesses	Debt Risk Score	Walk Score	Population	Average Income	Number of Cafes	Number of Parks/Playgrounds
0	5	Rosedale-Moore Park, Toronto	683	777	84	7672	213941	0.0	0.0
1	6	Bay Street Corridor, Toronto	4324	755	99	4787	40598	2.0	1.0

```
[196]: Preferred_Neighborhoods['Latitude']=''
Preferred_Neighborhoods['Longitude']=''
Preferred_Neighborhoods
```

	Cluster Label	Neighborhood	Businesses	Debt Risk Score	Walk Score	Population	Average Income	Number of Cafes	Number of Parks/Playgrounds	Latitude	Longitude
0	5	Rosedale-Moore Park, Toronto	683	777	84	7672	213941	0.0	0.0	0.0	43.6904
1	6	Bay Street Corridor, Toronto	4324	755	99	4787	40598	2.0	1.0	43.6645	-79.3872

```
[197]: i=0
while i<len(Preferred_Neighborhoods.index):
    j=0
    while j<len(df_combined):
        if Preferred_Neighborhoods.iloc[i,1]==df_combined.iloc[j,1]:
            Preferred_Neighborhoods.iloc[i,9]=df_combined.iloc[j,7]
            Preferred_Neighborhoods.iloc[i,10]=df_combined.iloc[j,8]
        j+=1
    i+=1
Preferred_Neighborhoods
```


	Cluster Label	Neighborhood	Businesses	Debt Risk Score	Walk Score	Population	Average Income	Number of Cafes	Number of Parks/Playgrounds	Latitude	Longitude
0	5	Rosedale-Moore Park, Toronto	683	777	84	7672	213941	0.0	0.0	43.6904	-79.3833
1	6	Bay Street Corridor, Toronto	4324	755	99	4787	40598	2.0	1.0	43.6645	-79.3872

Let's save Preferred_Neighborhood as a csv before putting a map together.

Charts below illustrate how the two neighborhoods match up head to head.


```
[9]: Image(filename='Selection_002.png')
```

[9]:




```
[10]: Image(filename='Selection_003.png')
```

[10]:



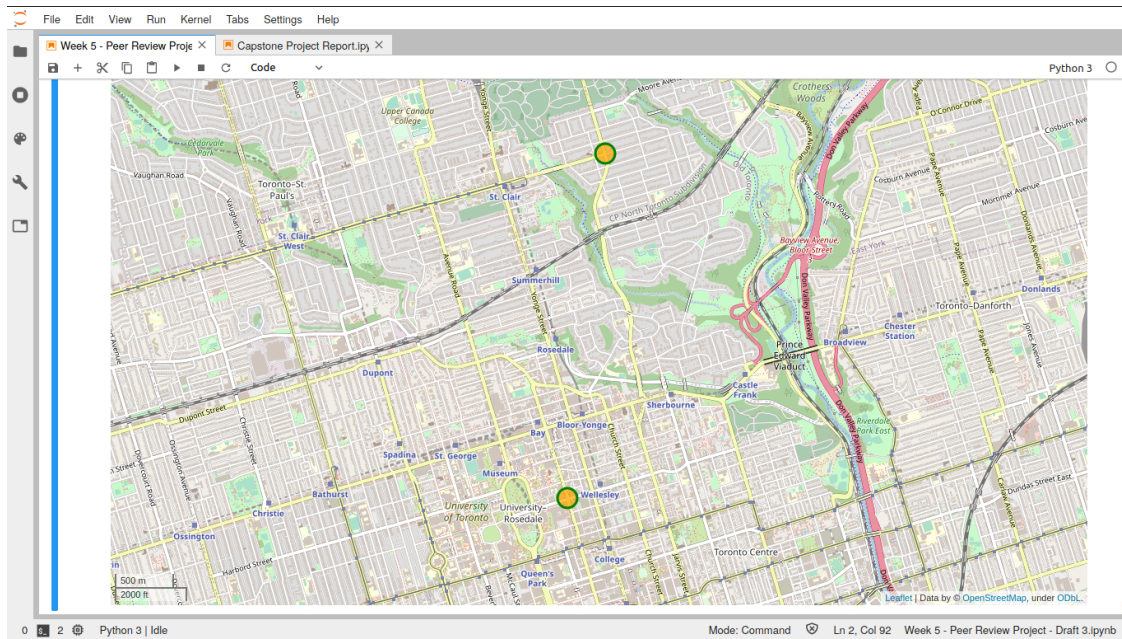
```
[11]: Image(filename='Selection_004.png')
```

[11]:



```
[12]: Image(filename='Screenshot from 2020-05-24 12-09-26.png')
```

[12]:



The proximity between the two neighborhoods is obvious (marked in yellow circles with green borders). They are located less than 5 miles of each other, an indication that they may share some of the traffic in the area.

1.6 Discussion

Caveats: Just a few caveats before going into the actual discussion part:

- i) I could not find complete information on all 140 census recognized neighborhoods in Toronto. Only 74 Toronto neighborhoods were ultimately included in the analysis.
- ii) This is just an academic exercise. I have never been to Toronto and my knowledge about the city is limited to the Google searches for this study. So some of my conclusions might seem obvious or ridiculous to someone familiar with the area.

Eight clusters were created by K mean clustering 74 Toronto neighborhoods based on seven desirable features. The two neighborhoods selected in the end are 'Bay Street Corridor' and 'Rosedale-Moore Park'. The two neighborhoods have numerous established and successful businesses. Residents of Rosedale-Moore Park are affluent and those in 'Bay Street Corridor' are very close to being top 25% of earners in Toronto. Both neighborhood have residents with high credit scores and are very walkable. As a bonus the two neighborhoods are within 5 km (~3.1 mi) of each other, likely sharing some traffic.

There are a few differences in the neighborhoods too. Bay Street Corridor is the financial district of Toronto and a major thoroughfare in the Downtown area.

Rosedale-Moore Park is an affluent neighborhood with lots of green space. It has been ranked as the best neighborhood in Toronto to live by Toronto Life.

What's interesting is in spite of the differences, both draw new faces from outside on top of the local residents. And both have more than enough financial strength to easily integrate a new business, provided it has a good concept which is executed well. This is where people making the business decisions need to take over. The ultimate choice will depend on their budget, past experiences and what they feel comfortable with as a company.

1.7 Conclusion

There are 140 census recognized neighborhoods in Toronto. The goal of the project was to create a short list of 1 to 5 neighborhoods where opening a new coffee shop would make good business sense. Cluster analysis (K means) was performed on 74 neighborhoods for which I could find all the necessary data on the factors discussed in the introduction section. The two neighborhoods finally chosen are: Bay Street Corridor and Rosedale-Moore Park. Bay Street Corridor is a busy thoroughfare in the financial district of Toronto and Rosedale-Moore Park is an affluent neighborhood with a lot of green space. In spite of their differences both are financially strong, have the ability to draw new people and are walkable; the perfect combination for a new coffee shop. People making the business decisions will decide which neighborhood to go with based on budget, experience and the style of business they are comfortable with.

1.8 Acknowledgment

The list of packages used in the various steps of this project includes numpy, scipy, pandas, request, BeautifulSoup, matplotlib, functools, folium and Scikit-Learn. Special thanks to all the contributors of these free open source packages for giving so much to the community.

1.9 References

1. <https://developer.foursquare.com/places>
2. <https://nominatim.org/>
3. https://en.wikipedia.org/wiki/Demographics_of_Toronto_neighbourhoods
4. https://open.toronto.ca/catalogue/?sort=last_refreshed%20desc

[]: